

Unsupervised learning using clustering— Mining the 20 Newsgroups Dataset

Haining Pan

September 18, 2022

1 Introduction

In this report, we study the similarity of newsgroups by utilizing several unsupervised learning clustering algorithms, including ‘Kmeans’, ‘Agglomerative clustering’, and ‘DBSCAN’. The newsgroups are composed of text files, such as posts, messages, and emails. We use the built-in dataset ‘fetch_20newsgroups’ in ‘Sklearn’. [1] The dataset carries the labels, therefore, our ultimate goal here is to see how well we can reproduce the given label using unsupervised learning clustering. This project is inspired by and generalized on [2].

2 Dataset

The 20 newsgroups dataset contains around 18000 posts on 20 topics. In this project, we choose 4 topics, ‘rec.sport.baseball’, ‘talk.politics.guns’, ‘comp.graphics’, ‘sci.med’, which contains 3867 posts (samples).

2.1 Preprocessing

2.2 Bag of word

We use the bag-of-word (BoW) model, which converts a text into a vector by counting the frequency of each word. Therefore, the number of features is just the number of different words in the whole dataset.

However, there are some words does not carry too much information, such as the name, numbers, and stop words. Therefore, we also remove them by using ‘NLTK’ package.

Furthermore, to account for the variant morphosis of the same word (e.g., ‘sit’ and ‘sitting’, ‘desk’ and ‘desks’), we ‘lemmatize’ words in order to create redundant features by using ‘WordNetLemmatizer’ in ‘NLTK’.

2.3 TF-IDF

Before we feed our vectors of term frequency into the clustering algorithms, we address the issue of the information that each word carries, because the more common a word appears in the whole sample, the less distinct information it carries. This is implemented by using the term frequency-inverse document frequency (TF-IDF). Therefore, the final data we feed into the clustering algorithms is the sparse matrix of the size of (n_samples, n_features) after the data preprocessing, where we name it ‘data’ in the code.

3 Clustering

In this section, we perform the clustering on the sparse matrix of ‘data’. To save the pages, we directly show results, while the code is provided in the GitHub repository. [3]

3.1 Kmeans

We start with the Kmeans, and sweep the number of clusters, as shown in Fig. 1, where we use Silhouette score and inertia as the metrics to evaluate the clustering results.

In Fig. 1, we notice the number of clusters $k = 4$ seems to be an inflection point in Elbow Method, while $k = 8$ gives the largest Sihouette score. Therefore, we treat these two ($k = 4$ and $k = 8$) as candidates, which will be later used to compare with the original labels by calculating the Rand Index.




Figure 1: Silhouette score (left axis) and inertia (right axis) as a function of the number of clusters in Kmeans.

3.2 Agglomerative clustering

We next apply the agglomerative clustering as shown in Fig. 2.




Figure 2: Silhouette score as a function of the number of clusters in Agglomerative clustering

In Fig. 2, we notice that $k = 6$ has the highest Silhouette score, while $k = 5$ is roughly the same. Therefore, we save these two estimators which are later used to compare with the original labels.

3.3 DBSCAN

Finally, we apply the DBSCAN and use ‘GridSearchCV’ to find the best ϵ and ‘min_samples’. We vary ϵ between 0.5 and 1.5 in a step of 0.1, and ‘min_samples’ between 2 and 15, and calculate the corresponding Silhouette score, as shown in Fig. 3.

In Fig. 3, we find most of the parameters give negative Silhouette scores, and only a few specific points give positive Silhouette scores, which are marked in red. Therefore, we will focus on these red dots, and see whether they can provide us with meaningful clusters in the next step.

We apply the DBSCAN again with parameters shown in red dots in Fig. 3. The results are shown in Table 1.

From Table 1, we notice that 1) most of the models have either just two clusters or unreasonably many clusters, which are not meaningful. Therefore, we doubt the applicability of DBSCAN here, since the original data may be highly inhomogeneous in the density. Nevertheless, we still choose models with 4, 5, and 6 clusters above, and see what the clusters are in those models.




Figure 3: Silhouette score as a function of ϵ and ‘min_samples’. The red dots mark the parameters with positive Silhouette scores.

Table 1: Silhouette score and the number of clusters for the DBSCAN model with different parameters of ϵ and ‘num_samples’.

(ϵ ,num_samples)	(0.5, 4)	(0.5, 5)	(0.5, 6)	(0.5, 7)	(0.6, 5)	(0.6, 6)	(0.6, 7)	(0.7, 6)
Silhouette score	0.0028	0.0028	0.0028	0.0028	0.0028	0.0028	0.0028	0.0028
Num_clusters	2	2	2	2	2	2	2	2
(ϵ ,num_samples)	(0.7, 7)	(0.8, 2)	(0.9, 2)	(0.9, 3)	(1, 2)	(1, 3)	(1.1, 2)	(1.1, 3)
Silhouette score	0.0028	0.0346	0.0731	0.0028	0.1008	0.0307	0.1013	0.0498
Num_clusters	2	423	564	175	651	273	619	314
(ϵ ,num_samples)	(1.1, 4)	(1.1, 5)	(1.2, 2)	(1.2, 3)	(1.2, 4)	(1.3, 3)	(1.3, 4)	(1.3, 5)
Silhouette score	0.0235	0.0070	0.0274	0.0071	0.0023	0.0004	0.0007	0.0003
Num_clusters	181	102	334	150	104	5	4	3
(ϵ ,num_samples)	(1.3, 6)	(1.3, 7)	(1.3, 8)	(1.3, 9)	(1.3, 10)			
Silhouette score	0.0020	0.0020	0.0020	0.0020	0.0020			
Num_clusters	2	2	2	2	2			

4 Final evaluation

Finally, we use the adjusted Rand Index to test the clustering from the different models above. For Kmeans, we obtain 0.5 for $k = 4$ and 0.34 for $k = 8$. For agglomerative clustering, we obtain 0.6 for $k = 5$ and 0.59 for $k = 6$. For DBSCAN, we obtain all very small values of Rand Index: 0.00026 for (1.3, 3), 0.00038 for (1.3, 4), and 0.00072 for (1.3,5).

To show explicitly what each cluster contains, we print 1. the size of each cluster; 2. the size of original labels that each cluster contains; 3. Top 10 words in each cluster.

For kmeans with $k = 4$, we have

```
cluster_0: 698 samples
Top Keywords: gun/wa/people/fbi/government/article/atf/right/law/batf
               talk.politics.guns: 697 samples
               sci.med: 1 samples
cluster_1: 751 samples
Top Keywords: game/team/wa/player/baseball/article/hit/win/pitcher/brave
               rec.sport.baseball: 750 samples
               talk.politics.guns: 1 samples
cluster_2: 670 samples
Top Keywords: graphic/image/file/program/format/need/university/looking/know/bit
               comp.graphics: 660 samples
               sci.med: 8 samples
               rec.sport.baseball: 2 samples
cluster_3: 1748 samples
Top Keywords: wa/article/university/ha/know/like/just/doe/medical/computer
               sci.med: 981 samples
               comp.graphics: 313 samples
               rec.sport.baseball: 242 samples
               talk.politics.guns: 212 samples
```

For kmeans with $k = 8$, we have

cluster_0: 382 samples
Top Keywords: wa/fbi/atf/batf/government/waco/burn/dividian/article/people
talk.politics.guns: 381 samples
sci.med: 1 samples

cluster_1: 706 samples
Top Keywords: game/team/wa/player/baseball/hit/article/pitcher/brave/win
rec.sport.baseball: 705 samples
sci.med: 1 samples

cluster_2: 1435 samples
Top Keywords: university/wa/article/know/doe/just/ha/like/thanks/computer
comp.graphics: 537 samples
sci.med: 402 samples
rec.sport.baseball: 289 samples
talk.politics.guns: 207 samples

cluster_3: 317 samples
Top Keywords: gun/people/firearm/right/law/weapon/wa/like/article/criminal
talk.politics.guns: 317 samples

cluster_4: 439 samples
Top Keywords: image/graphic/file/program/format/bit/color/need/computer/display
comp.graphics: 435 samples
sci.med: 4 samples

cluster_5: 58 samples
Top Keywords: msg/food/sensitivity/chinese/restaurant/reaction/eat/allergic/people/effect
sci.med: 58 samples

cluster_6: 455 samples
Top Keywords: doctor/ha/wa/medical/patient/disease/pain/yeast/article/drug
sci.med: 449 samples
talk.politics.guns: 5 samples
comp.graphics: 1 samples

cluster_7: 75 samples
Top Keywords: shameful/chastity/surrender/gordon/bank/pittsburgh/science/computer/article/lyme
sci.med: 75 samples

For agglomerative clustering with 5 clusters, we have

cluster_0: 839 samples
Top Keywords: game/wa/team/player/baseball/article/hit/ha/win/think
rec.sport.baseball: 815 samples
sci.med: 13 samples
comp.graphics: 6 samples
talk.politics.guns: 5 samples

cluster_1: 791 samples
Top Keywords: gun/wa/people/fbi/article/government/right/atf/like/law
talk.politics.guns: 772 samples
comp.graphics: 9 samples
sci.med: 6 samples
rec.sport.baseball: 4 samples

cluster_2: 1416 samples
Top Keywords: wa/article/university/ha/know/medical/science/like/doe/computer
sci.med: 896 samples
comp.graphics: 227 samples
rec.sport.baseball: 166 samples
talk.politics.guns: 127 samples

cluster_3: 58 samples
Top Keywords: msg/food/sensitivity/chinese/restaurant/allergic/eat/reaction/people/flavor
sci.med: 58 samples

cluster_4: 763 samples
Top Keywords: graphic/image/file/program/university/need/format/computer/know/bit
comp.graphics: 731 samples
sci.med: 17 samples
rec.sport.baseball: 9 samples
talk.politics.guns: 6 samples

For agglomerative clustering with 6 clusters, we have

cluster_0: 791 samples
Top Keywords: gun/wa/people/fbi/article/government/right/atf/like/law
talk.politics.guns: 772 samples
comp.graphics: 9 samples
sci.med: 6 samples
rec.sport.baseball: 4 samples

cluster_1: 763 samples
Top Keywords: graphic/image/file/program/university/need/format/computer/know/bit
comp.graphics: 731 samples
sci.med: 17 samples
rec.sport.baseball: 9 samples
talk.politics.guns: 6 samples

cluster_2: 1416 samples
Top Keywords: wa/article/university/ha/know/medical/science/like/doe/computer
sci.med: 896 samples
comp.graphics: 227 samples
rec.sport.baseball: 166 samples
talk.politics.guns: 127 samples

cluster_3: 58 samples
Top Keywords: msg/food/sensitivity/chinese/restaurant/allergic/eat/reaction/people/flavor
sci.med: 58 samples

cluster_4: 803 samples
Top Keywords: game/team/wa/player/article/hit/baseball/ha/win/brave
rec.sport.baseball: 779 samples
sci.med: 13 samples
comp.graphics: 6 samples
talk.politics.guns: 5 samples

cluster_5: 36 samples
Top Keywords: jewish/baseball/come/hank/sandy/racking/lame/john/wa/article
rec.sport.baseball: 36 samples

For DBSCAN with (1.3,3), we have

cluster_0: 3754 samples
Top Keywords: wa/ha/world/like/university/heard/research/article/similar/japanese
rec.sport.baseball: 984 samples
comp.graphics: 953 samples
sci.med: 930 samples
talk.politics.guns: 887 samples

cluster_1: 4 samples
Top Keywords: wa/article/university/ha/like/know/just/think/people/gun
sci.med: 4 samples

cluster_2: 5 samples
Top Keywords: leung/drink/ng/ho/lactose/intolerance/experiencing/milk/ton/hardly
sci.med: 5 samples

cluster_3: 3 samples
Top Keywords: centipede/millipede/posionous/pes/bitten/liable/clarification/know/concerted/soon
comp.graphics: 3 samples

cluster_4: 3 samples
Top Keywords: playmation/retail/ren/hitchhiking/kenbaer/anjon/hash/sell/scn/marino
sci.med: 3 samples

For DBSCAN with (1.3,4), we have

cluster_0: 3725 samples
Top Keywords: wa/ha/world/like/university/article/research/playmation/know/information
rec.sport.baseball: 979 samples
comp.graphics: 948 samples
sci.med: 918 samples
talk.politics.guns: 880 samples

cluster_1: 4 samples
Top Keywords: wa/article/university/ha/like/know/just/think/people/gun
sci.med: 4 samples

cluster_2: 4 samples

```
Top Keywords: leung/drink/ng/ho/lactose/intolerance/experiencing/milk/ton/hardly
sci.med: 4 samples
cluster_3: 5 samples
Top Keywords: acne/oily/face/teenage/skin/wash/son/clearasil/dalacin/nose
sci.med: 5 samples
```

For DBSCAN with (1.3,5), we have

```
cluster_0: 3698 samples
Top Keywords: wa/ha/university/article/world/like/information/think/know/leung
rec.sport.baseball: 976 samples
comp.graphics: 945 samples
sci.med: 902 samples
talk.politics.guns: 875 samples
cluster_1: 5 samples
Top Keywords: wa/article/university/ha/like/know/just/think/people/gun
sci.med: 5 samples
cluster_2: 4 samples
Top Keywords: centipede/millipede/posionous/pes/bitten/liable/clarification/know/concerted/s
sci.med: 4 samples
```

5 Conclusion

In summary, the best performance of clustering is reached by agglomerative clustering with the Rand Index of ~ 0.6 . The explicit number of clusters does not matter that much between 5 and 6, as they basically refer to the same types of clustering. For the 6 clusters, there are two clusters (cluster3 and cluster5) that are not well distinguished, however, these two clusters have very small sizes, which can be reasonably ignored. Besides that, agglomerative clustering identifies most clusters correctly.

For kmeans, the best model is the one with 4 clusters although the model with 8 clusters has a larger Silhouette score. However, it is generally worse than agglomerative clustering, as it always leaves a cluster containing all four original labels, which indicates there are some texts that kmeans cannot distinguish. (For example, cluster3 in $k=4$) However, besides that cluster, all other three clusters are classified quite well, which is as good as the agglomerative method.

For DBSCAN, it cannot produce meaningful clusters, as it simply classifies all samples into one big cluster (which is the boundary with label -1 by further check) This could be due to the intrinsic inhomogeneous density in the dataset, therefore, DBSCAN is not very useful in such a text-based situation.

6 Reference

References

- [1] https://scikit-learn.org/stable/datasets/real_world.html#newsgroups-dataset
- [2] Chapter 3, Python Machine Learning By Example, Liu Yuxi, PacktPub
- [3] <https://github.com/hainingpan/UnsupervisedLearning/blob/master/UL.ipynb>